



UC Berkeley
Teaching Professor
Lisa Yan

CS61C

Great Ideas
in
Computer Architecture
(a.k.a. Machine Structures)



UC Berkeley
Lecturer
Justin Yokota

Caches III: Direct Mapped

- **EECS/CS Undergraduate Town Hall**,
Thursday April 11 4-5pm,
Wozniak Lounge
(tomorrow!)
• Data4All Kickoff event!
RSVP by today. EdStem
post #824 [\[link\]](#)

Warming up the Fully Associative Cache

[review]

- Suppose the cache starts **cold**.

- Load byte 0x43F 0x10F, 0x3
 - Load byte 0x5E2 0x178, 0x2
 - Load word 0x824 0x209, 0x0
 - Load word 0x5E0 0x178, 0x0
-

Val id	Tag	Data			
		11	10	01	00
1	0x10F
1	0x178
1	0x209
0

- Cache hit!** Tag 0x178 is valid!
- Read byte at 0x0 **offset** and return to processor.

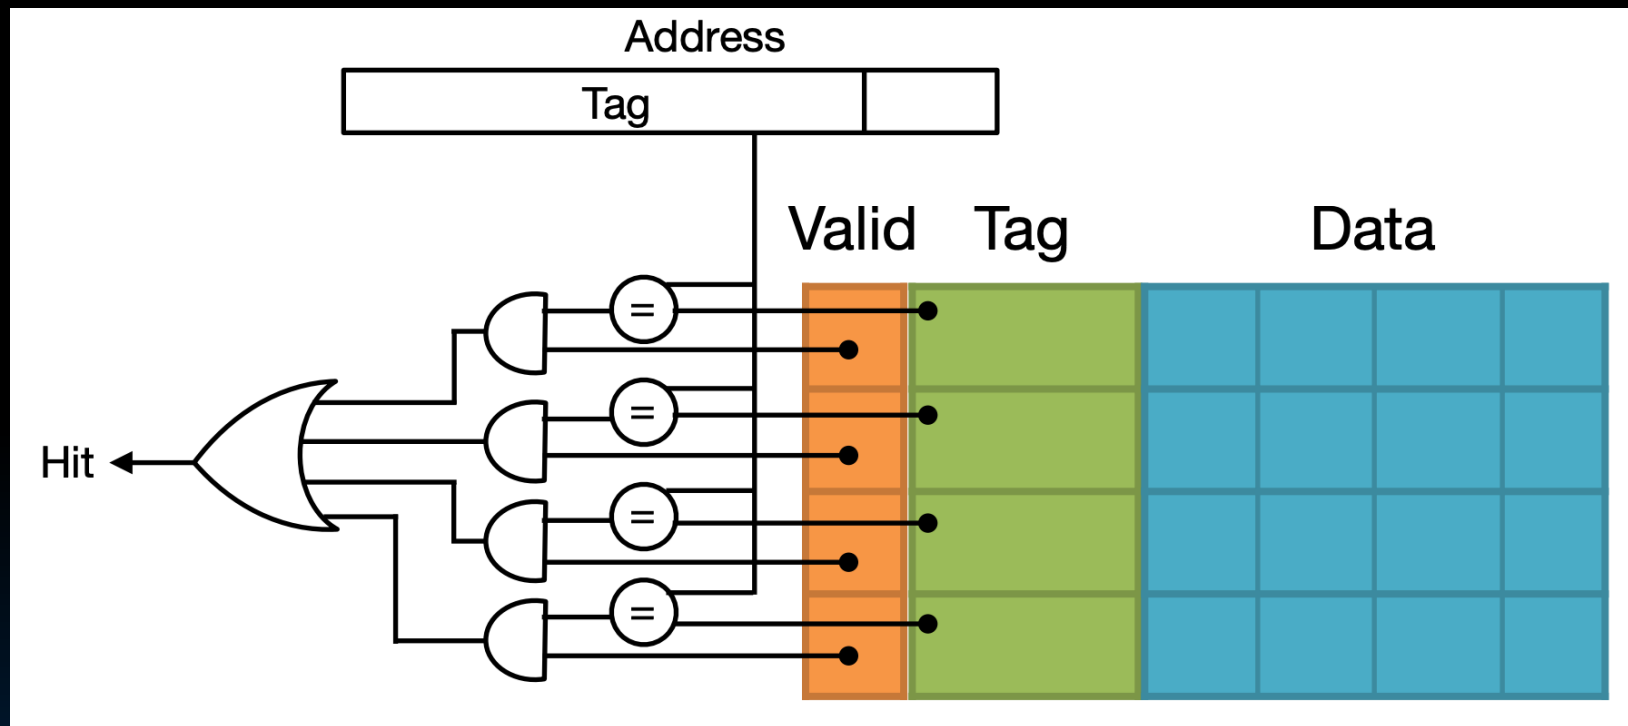
Memory address:	11	2	1 0
	(10b)		(2b)
	tag		offset

No access to main memory occurs on a **cache hit**.

- Cache line (i.e., cache) block
 - The smallest unit of memory that can be transferred between the main memory and the cache.
 - Each line has its own entry in the cache.
 - Line size/block size: The number of bytes in each cache line.
- When we bring data from the main memory into the cache, it is done in the granularity of a cache line (or cache block).
 - Typically **cache lines are 64 bytes**.
 - Cache blocks (cache lines) helps us take advantage of **spatial locality**.
- **Capacity:**
 - The total number of data bytes that can be stored in a cache.
 - For fully associative cache, **capacity = # lines * line size**.

Hardware: Fully Associative Cache

- A Fully Associative Cache must compare tags in parallel:




Need hardware comparator for every single entry!
 Sometimes infeasible; e.g., 1MB cache w/ 16B entries
 $\rightarrow 2^{20}/2^4 = 2^{16} = 16\text{K comparators.}$

Block Replacement Policies

- Block Replacement Policies
- Write Policies
- Direct Mapped Cache
- Types of Misses

A Warmed up Cache Still Can Miss

- Suppose the cache starts **cold**.
 - Load byte 0x43F 0x10F, 0x3
 - Load byte 0x5E2 0x178, 0x2
 - Load word 0x824 0x209, 0x0
 - Load word 0x5E0 0x178, 0x0
 - Load word 0x524 0x149, 0x0


 Fills up the cache

Val id	Tag	Data			
		11	10	01	00
1	0x10F
1	0x178
1	0x209
1	0x149



A Warmed up Cache Still Can Miss

By the end of instr 5, cache is now **warm**.

- 1. Load byte 0x43F 0x10F, 0x3
- 2. Load byte 0x5E2 0x178, 0x2
- 3. Load word 0x824 0x209, 0x0
- 4. Load word 0x5E0 0x178, 0x0
- 5. Load word 0x524 0x149, 0x0
- 6. Load byte 0x972



Cache miss!

Val id	Tag	Data			
		11	10	01	00
1	0x10F
1	0x178
1	0x209
1	0x149

When the FA cache **reaches capacity**, we can still miss.
Therefore we also need a policy for **block replacement!**

Block Replacement Policies (i.e., Eviction)

- Least Recently Used (LRU)
 - Replace the entry that has not been used for the longest time, i.e., has the oldest previous access.
 - Pro: **temporal locality!**
 - recent past use implies likely future use
 - This is a very effective policy
 - Con: Complicated hardware to keep track of access history → performance hit

Fully Associative Cache with LRU policy

By the end of instr 5, cache is now **warm**.

1. Load byte 0x43F

0x10F, 0x3
2. Load byte 0x5E2

0x178, 0x2
3. Load word 0x824

0x209, 0x0
4. Load word 0x5E0

0x178, 0x0
5. Load word 0x524

0x149, 0x0
6. Load byte 0x972

1001 0111 0010

1001 0111 0010

tag

0x25C

offset

0x2

Suppose that LRU = 0 means **most** recently used, and 3 means **least** recently used. After the end of instruction 5, what are the LRU tags on each row?

Val id	LRU	Tag	Data			
			11	10	01	00
1		0x10F
1		0x178
1		0x209
1		0x149

A.

LRU

0

1

2

3

B.

LRU

3

2

1

0

C.

LRU

3

1

2

0

D.

LRU

0

2

3

1

E. Something else/don't know

At the end of instruction 5, what are the LRU tags on each row?

CS 61C

Fully Associative Cache with LRU policy

▪ By the end of instr 5, cache is now warm.

1. Load byte 0x43F 0x10F, 0x3

2. Load byte 0x5E2 0x178, 0x2

3. Load word 0x824 0x209, 0x0

4. Load word 0x5E0 0x178, 0x0

5. Load word 0x524 0x149, 0x0

6. Load byte 0x972 1001 0111 0010

tag offset

0x25C 0x2

Suppose that LRU = 0 means most recently used, and 3 means least recently used.

After the end of instruction 5, what are the LRU tags on each row?

Val id	LRU	Tag	11	10	01	00
1		0x10F
1		0x178
1		0x209
1		0x149

A. LRU

B. LRU

C. LRU

D. LRU

0	3	3	0
1	2	1	2
2	1	2	3
3	0	0	1

E. Something else/don't know

A

0%

B

0%

C

0%

D

0%

Something else/don't know

0%

Fully Associative Cache with LRU policy

By the end of instr 5, cache is now **warm**.

- 1. Load byte 0x43F 0x10F, 0x3
- 2. Load byte 0x5E2 0x178, 0x2
- 3. Load word 0x824 0x209, 0x0
- 4. Load word 0x5E0 0x178, 0x0
- 5. Load word 0x524 0x149, 0x0
- 6. Load byte 0x972 1001 0111 0010

Val id	LRU	Tag	Data			
			11	10	01	00
1	3	0x10F
1	1	0x178
1	2	0x209
1	0	0x149



Suppose that LRU = 0 means **most** recently used, and 3 means **least** recently used.
 After the end of instruction 5,
 what are the LRU tags on each row?

Fully Associative Cache with LRU policy

By the end of instr 5, cache is now **warm**.

- Load byte 0x43F 0x10F, 0x3
- Load byte 0x5E2 0x178, 0x2
- Load word 0x824 0x209, 0x0
- Load word 0x5E0 0x178, 0x0
- Load word 0x524 0x149, 0x0
- Load byte 0x972 0x25C, 0x2

Val id	LRU	Tag	Data			
			11	10	01	00
1	0	0x25C
1	2	0x178
1	3	0x209
1	1	0x149

a. Cache miss! Tag 0x25C is not valid.

b. **Replace according to block replacement policy.**

Load into cache the **4-byte block** from 0x970 to 0x973. Mark **valid bit**. Update block replacement policy fields.

c. Read byte at **0x2 offset** and return to processor.

Suppose our cache uses LRU (Least Recently Used).

Block Replacement Policies (i.e., Eviction)

- Least Recently Used (LRU)
 - Replace the entry that has not been used for the longest time, i.e., has the oldest previous access.
 - Pro: **temporal locality**!
 - recent past use implies likely future use
 - This is a very effective policy
 - Con: Complicated hardware to keep track of access history → performance hit
- Most Recently Used (MRU)
 - Replace the entry that has the newest previous access.
- First In, First Out (FIFO)
 - Replace the oldest block in the set (queue).
- Last In, First Out (LIFO)
 - Replace the newest block in the set (stack).
- Random
 - Works surprisingly okay (when given a low temporal locality workload)

LRU is ideal for **temporal locality**, but the last three are used more commonly in practice.

Reasonable approximations to LRU and MRU without adding too much excess hardware.

Note both ignore order of accesses after (before) the first (last) access.



Write Policies

- Block Replacement Policies
- Write Policies
- Direct Mapped Cache
- Types of Misses



What about Stores?

By the end of instr 5, cache is now **warm**.

- 1. Load byte 0x43F 0x10F, 0x3
- 2. Load byte 0x5E2 0x178, 0x2
- 3. Load word 0x824 0x209, 0x0
- 4. Load word 0x5E0 0x178, 0x0
- 5. Load word 0x524 0x149, 0x0
- 6. Load byte 0x972 0x25C, 0x2
- 7. Store byte 0x524 0x149, 0x0

Val id	LRU	Tag	Data			
			11	10	01	00
1	0	0x25C
1	2	0x178
1	3	0x209
1	1	0x149

Cache hit!...and then?

How do we handle stores?

Write-through vs. Write-back Policies

- **Store** instructions write to memory, which **changes values**.
- Hardware needs to ensure that cache and memory have consistent information.
- **Write-through:**
 - Write to the cache and memory at the same time.
 - (writes to memory take longer)

Write-through vs. Write-back Policies

[LRU column fixed post-class]

- Store instructions write to memory, which **changes values**.
- Hardware needs to ensure that cache and memory have consistent information.
- Write-through:**
 - Write to the cache and memory at the same time.
 - (writes to memory take longer)
- Write-back:**
 - Write data in cache and set a **dirty bit** to 1.
 - When this block gets evicted from the cache (and “back” to memory), write to memory.

Val id	Dir ty	LRU	Tag	Data			
				11	10	01	00
1	0	1	0x25C
1	0	2	0x178
1	0	3	0x209
1	1	0	0x149

Store byte 0x524 0x149, 0x0

Cache hit w/write-back:
 update data **within cache block**, and only write back to memory when this block is evicted.

Write-back policies require an additional **dirty bit** field.

Fully Associative Cache w/ LRU, Write-Back

Note: LRU also updates for all blocks on writes as well!! Because you are “using” blocks.
[LRU column fixed post-class]

Val id	Dir ty	LRU	Tag	Data			
				11	10	01	00
1	0	1	0x25C
1	0	2	0x178
1	0	3	0x209
1	1	0	0x149

Store byte 0x524

0x149, 0x0

Write-back:

- Write data in cache and set a **dirty bit** to 1.
- When this block gets evicted from the cache (and “back” to memory), write to memory.

Cache hit w/write-back:

update data **within cache block**, and only write back to memory when this block is evicted.

Write-through vs. Write-back Policies

- Store instructions write to memory, which **changes values**.
- Hardware needs to ensure that cache and memory have consistent information.

- **Write-through:**

- Write to the cache and memory at the same time.
- (writes to memory take longer)

Very simple to
implement

- **Write-back:**

- Write data in cache and set a **dirty bit** to 1.
- When this block gets evicted from the cache (and “back” to memory), write to memory.

(typically) lower traffic to
memory, because you likely
write multiple times before
evicting from cache

Direct Mapped Cache

- Block Replacement Policies
- Write Policies
- Direct Mapped Cache
- Types of Misses

Cache Design: Placement Policies

Fully
Associative
Cache

Put a new
block
anywhere

Set-Associative
Cache

(out of scope)

**Direct
Mapped
Cache**

Put a new block in
one specific place

(up next)

Fully associative caches need expensive hardware.

Direct Mapped Cache

- **Placement policy:** Each memory address is associated with exactly one possible block in the cache.
 - To check for existence in the cache, we only need to look in a single location in the cache.

Val id	Dir ty	Tag	Data			
			11	10	01	00
...
...
...
...

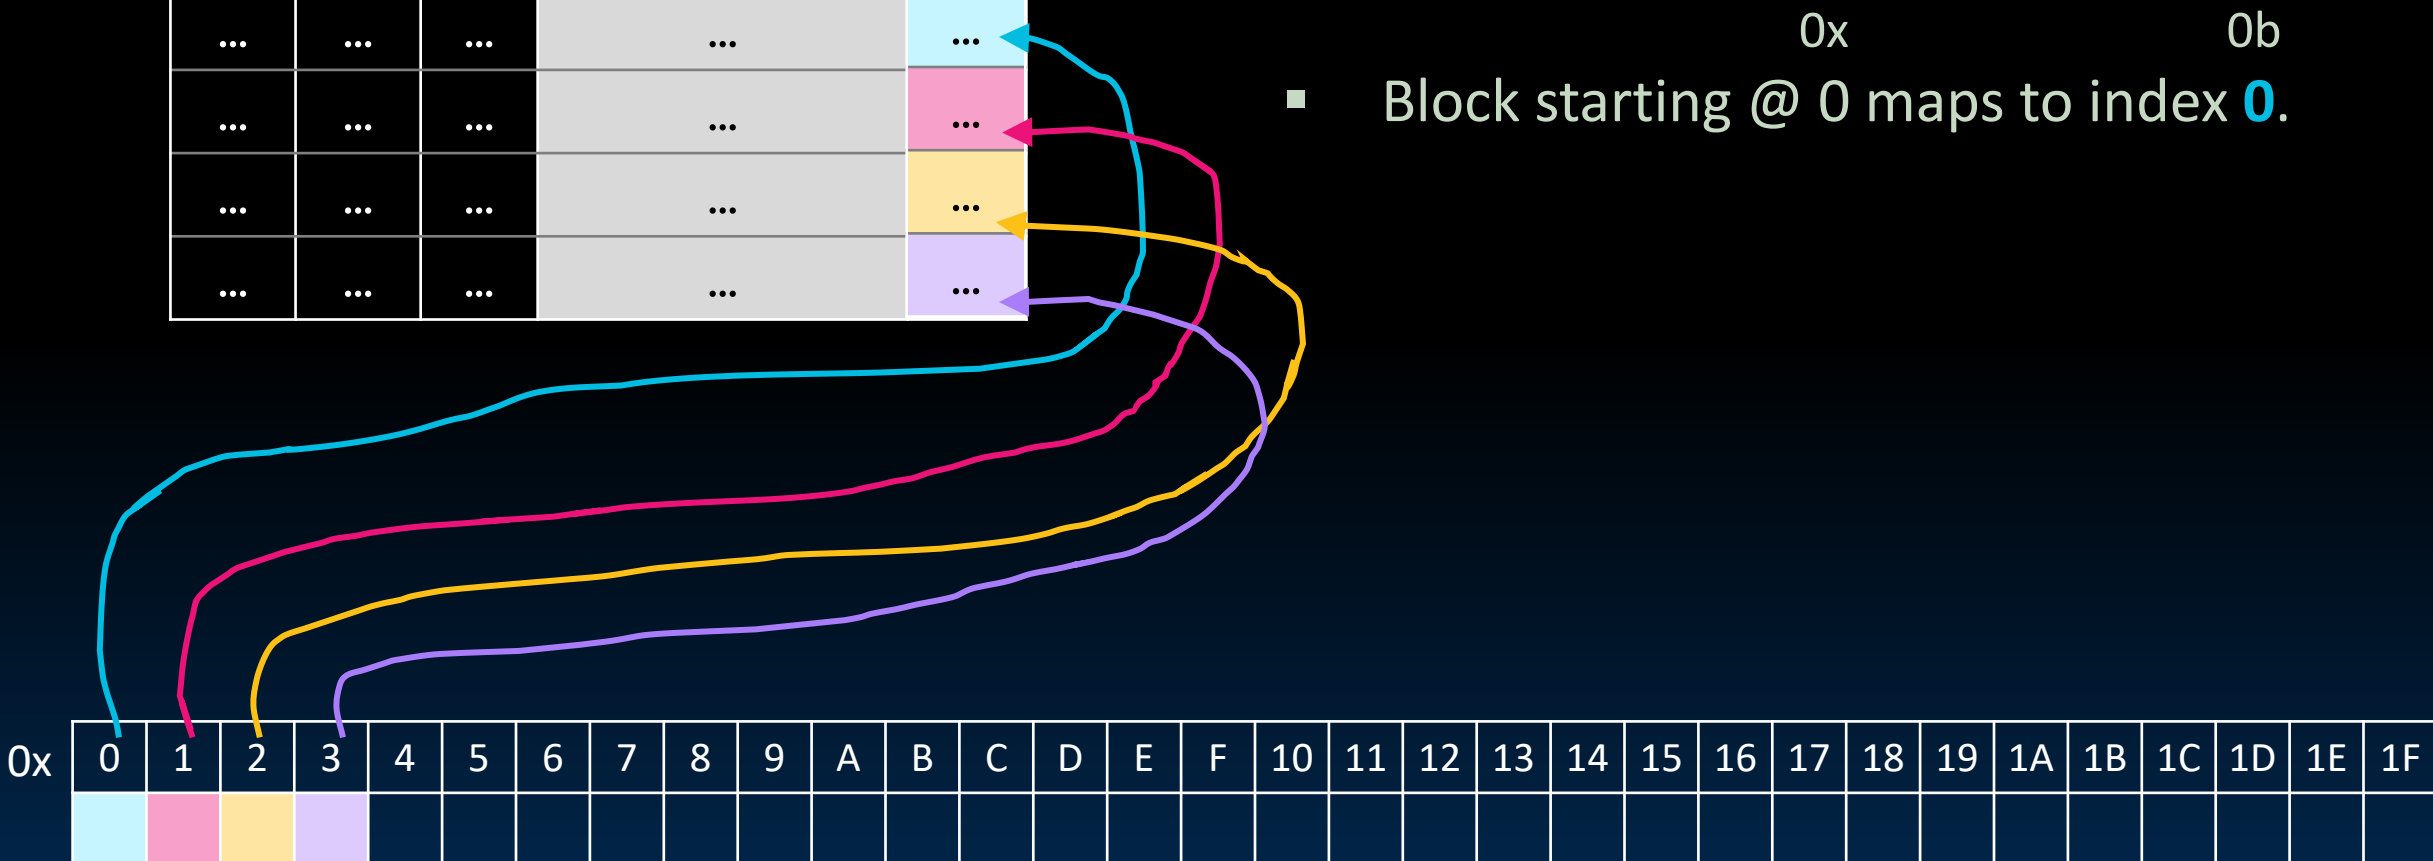
How do we ensure this?

(1/2) The Idea behind Direct Mapped Caches

4B Direct Mapped Cache
(block size 1B)

Val id	Dir ty	LRU	Tag	Data
...
...
...
...

- Block starting @ 0 maps to index **0**.



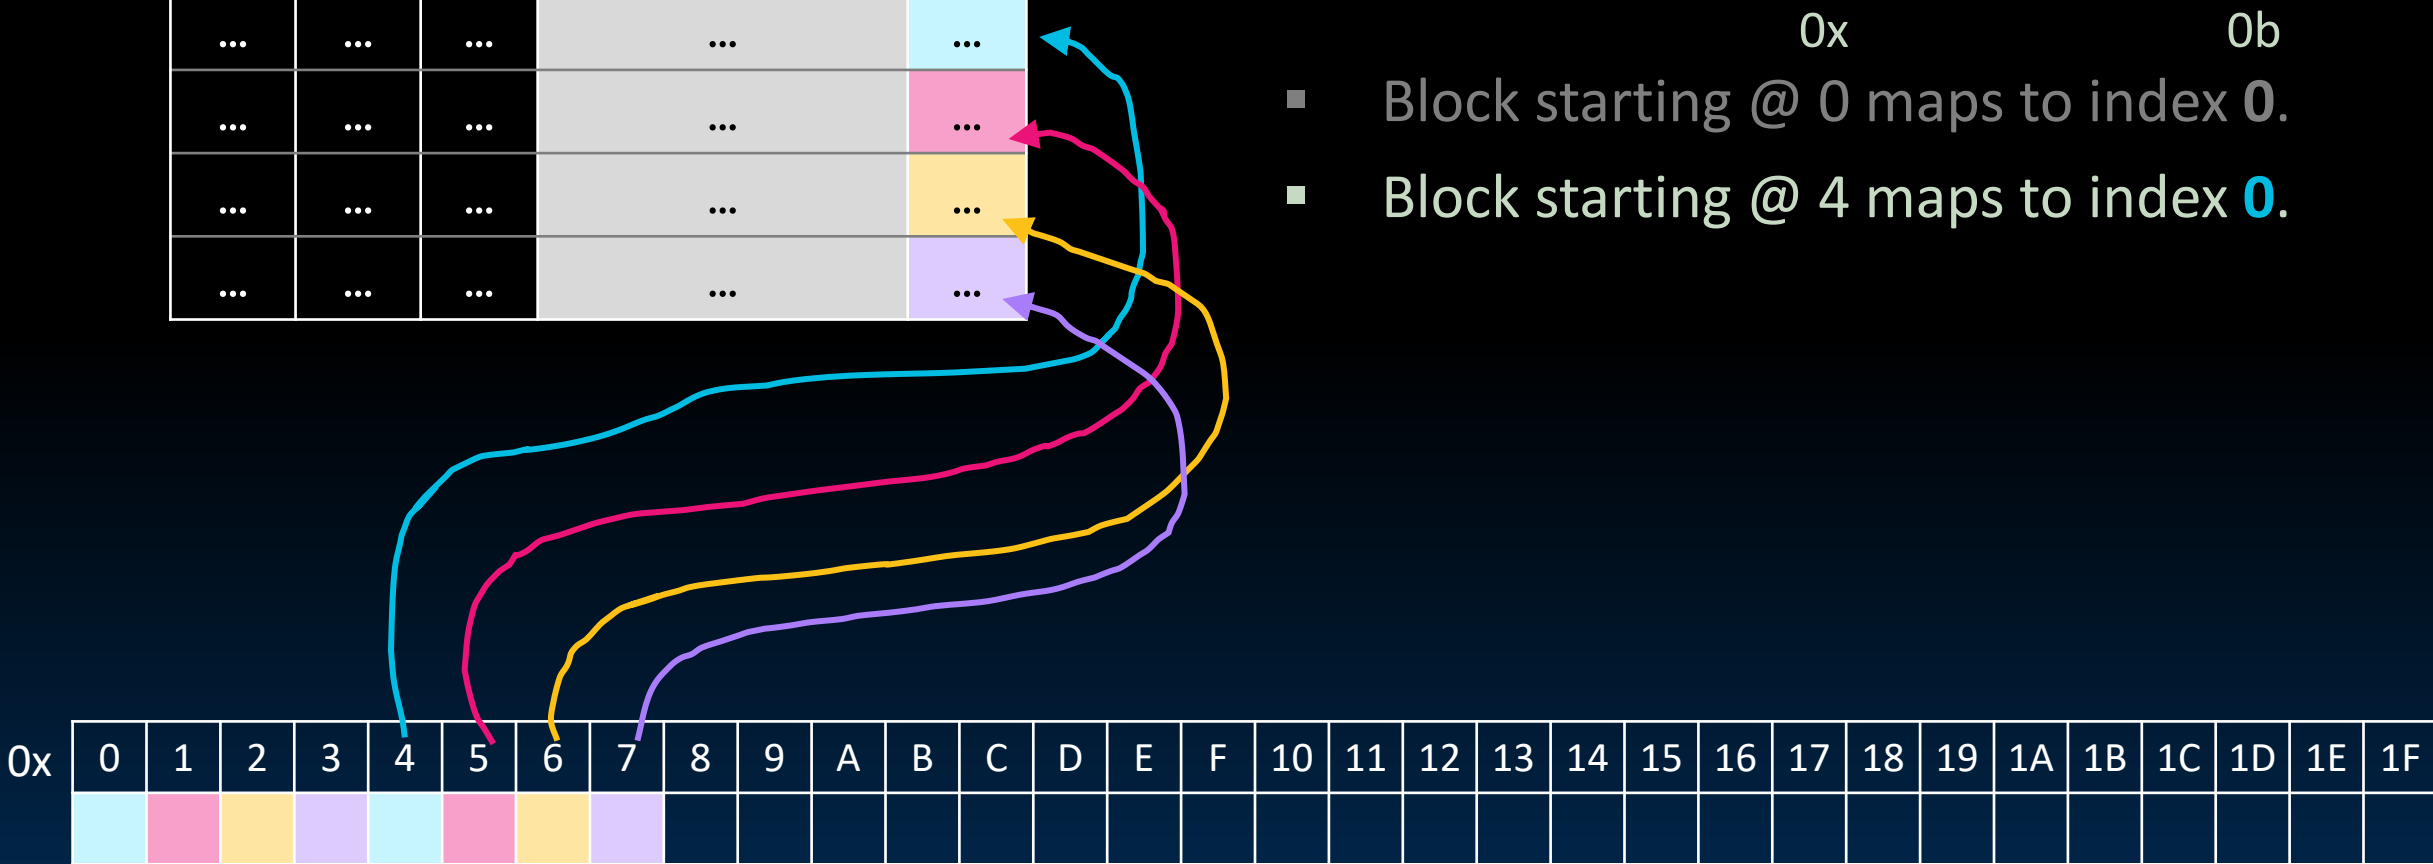
Memory

(1/2) The Idea behind Direct Mapped Caches

4B Direct Mapped Cache
(block size 1B)

Val id	Dir ty	LRU	Tag	Data
...
...
...
...

- Block starting @ 0 maps to index **0**.
- Block starting @ 4 maps to index **0**.



Memory

(1/2) The Idea behind Direct Mapped Caches

Valid	Dirty	LRU	Tag	Data
...
...
...
...

4B Direct Mapped Cache (block size 1B)

- Block starting @ 0 maps to index **0**.
- Block starting @ 4 maps to index **0**.
- Block starting @ 8 maps to index **0**.
- Block starting @ C maps to index **0**.
- etc.

Different addresses, get **same block index** but different **tags**.


0x	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F

Memory

(2/2) The Idea behind Direct Mapped Caches

Val id	Dir ty	LRU	Tag	Data	
				01	00
...
...
...
...

8B Direct Mapped Cache
(block size 2B)




0x	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F

Memory

(2/2) The Idea behind Direct Mapped Caches

Val id	Dir ty	LRU	Tag	Data	
				01	00
...
...		
...		
...		

8B Direct Mapped Cache
(block size 2B)



- Block starting @ 0 maps to index **0**.
- Block starting @ 8 maps to index **0**.
- Block starting @ 10 maps to index **0**.
- Block starting @ 18 maps to index **0**.
- etc.

Different addresses, get **same block index** but different **tags**.

0x	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F

Memory

(2/2) The Idea behind Direct Mapped Caches

8B Direct Mapped Cache
(block size 2B)

Val id	Dir ty	LRU	Tag	Data	
				01	00
...
...
...
...

0x	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F

Memory



Direct Mapped Cache

- **Placement policy:** Each memory address is associated with exactly one possible block in the cache.
 - To check for existence in the cache, we only need to look in a single location in the cache.

Val id	Dir ty	Tag	Data			
			11	10	01	00
...
...
...
...



tag to check if
have correct block

index to
select block


byte offset
within block

In direct mapped caches, split address to also get **index** of cache block to directly map to.

(2/2) The Idea behind Direct Mapped Caches

Val id	Dir ty	LRU	Tag	Data	
				01	00
...		
...		
...		
...		

8B Direct Mapped Cache
(block size **2B**)



- Block starting @ 0 maps to index **0**. Tag 00
- Block starting @ 8 maps to index **0**. Tag 01
- Block starting @ 10 maps to index **0**. Tag 10
- Block starting @ 18 maps to index **0**. Tag **11**

Different addresses, get **same block index** but different **tags**.

1 1**00**x

0x	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F

Memory

Fill in the blank

Val id	Dir ty	Tag	Data			
			11	10	01	00
...
...
...
...

[activity
slide]

Suppose we have the following **direct mapped** cache, for **12b** addresses.

	A.	B.	C.	D.	E.	F.
1. What is the block size / line size, in bytes?	2	4	8	10	16	Other
2. What is the capacity, in bytes?	2	4	8	10	16	Other
3. For a 12b address, how many bits is the byte offset?	2	4	8	10	16	Other
4. For a 12b address, how many bits is the index ?	2	4	8	10	16	Other
5. For a 12b address, how many bits is the tag?	2	4	8	10	16	Other

Put a pin on the answers for **direct mapped caches.**

0



Fill in the blank

Suppose we have the following **direct mapped** cache, for **12b** addresses. **A.**

- | | | | | | | |
|---|---|---|---|----|----|-------|
| 1. What is the block size / line size, in bytes? | 2 | 4 | 8 | 10 | 16 | Other |
| 2. What is the capacity, in bytes? | 2 | 4 | 8 | 10 | 16 | Other |
| 3. For a 12b address, how many bits is the byte offset? | 2 | 4 | 8 | 10 | 16 | Other |
| 4. For a 12b address, how many bits is the index ? | 2 | 4 | 8 | 10 | 16 | Other |
| 5. For a 12b address, how many bits is the tag? | 2 | 4 | 8 | 10 | 16 | Other |

Val id	Dir ty	LRU	Tag	11	10	01	00
...
...
...
...

[activity
slide]



Terminology, Part 1 (for 12b addresses!)

Cache block/line: A single entry in the cache

1. Block size / line size: # bytes per cache block
4 bytes

2. Capacity 4 x 4 bytes = **16 bytes**

▪ Total # data bytes that can be stored in a cache

3. Offset $\log_2(\text{block size}) = \mathbf{2 \text{ bits}}$

▪ Identifies byte offset of data stored within a given cache block

4. Index $\log_2(\# \text{ lines}) = \mathbf{2 \text{ bits}}$

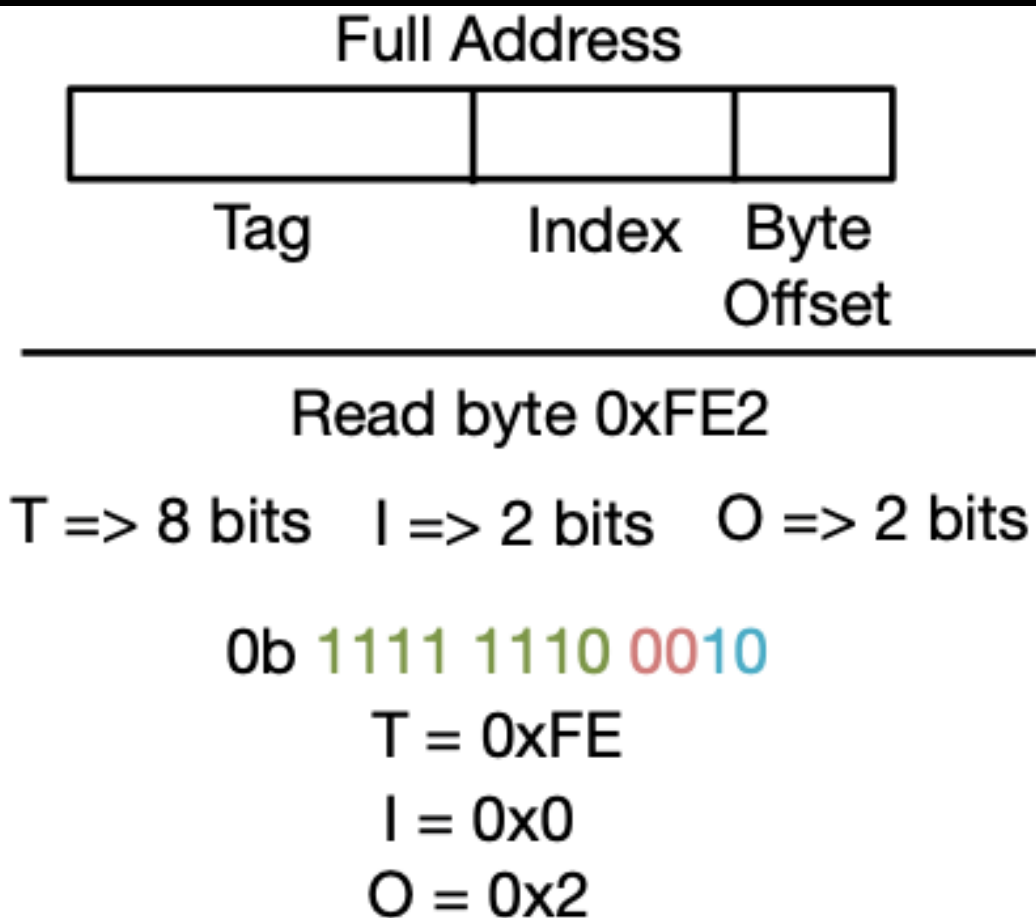
▪ Selects block

5. Tag # address bits - # offset bits - # index bits = **8 bits**

▪ Identifies (checks) if given cache block is correct

Val id	Tag	Data			
		11	10	01	00

Direct Mapped Example

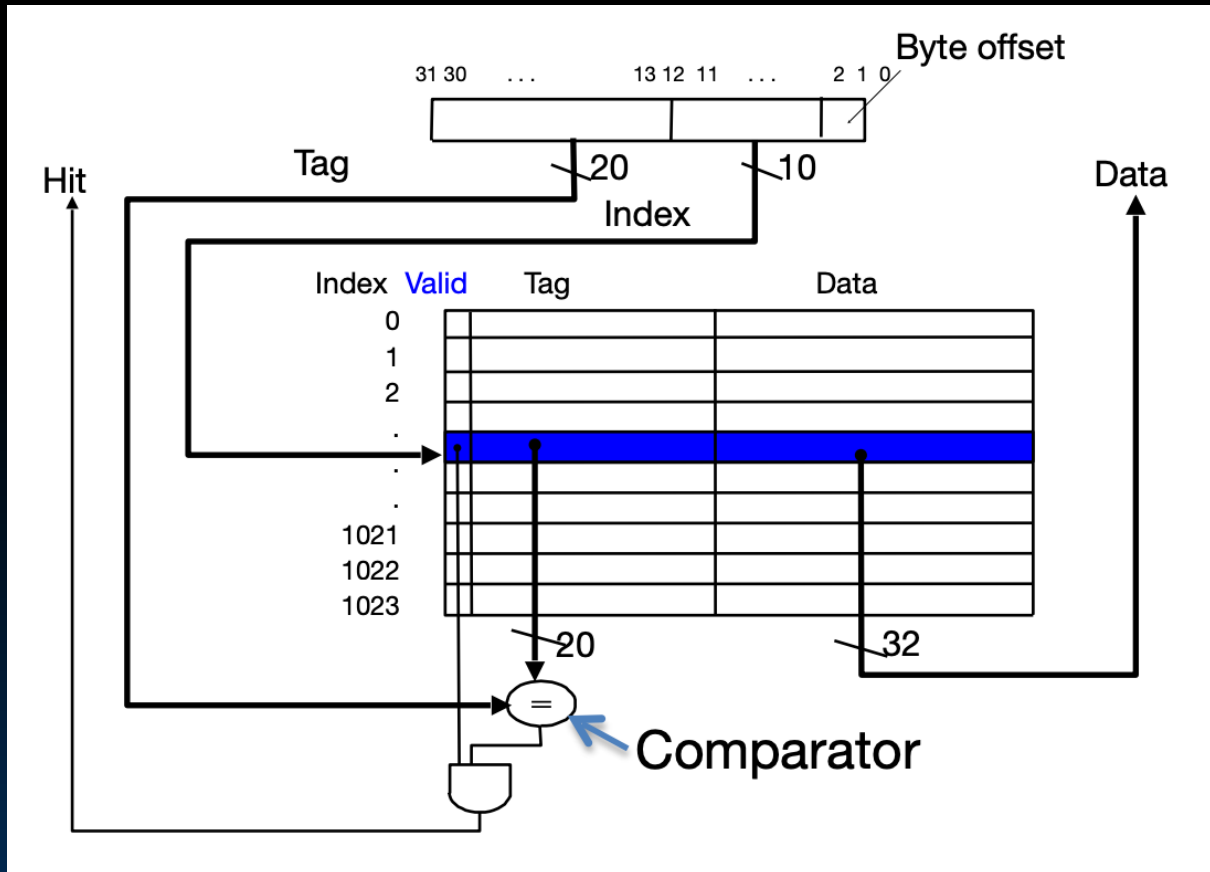


	Valid	Dirty	Tag	Data			
				11	10	01	00
0	1	0	0xFE
1	0
2	0
3	0

4 Bytes

Cache Miss => bring in entire line

Hardware: Direct Mapped Cache



4B blocks, 4KB data

Much simpler than Fully Associative! Just check one tag.

Types of Misses

- Block Replacement Policies
- Write Policies
- Direct Mapped Cache
- Types of Misses

Direct Mapped: What is Possible?

- What policies can be implemented for a **direct-mapped cache**?
Select all that apply.

1. Block Replacement (Eviction) Policy

- A. Least Recently Used
- B. Most Recently Used
- C. FIFO
- D. Random
- E. None of the above

2. Write-Back Policy

- A. Write-through
(memory access per write)
- B. Write-back
(dirty bit and write on eviction)

What policies can be implemented for a direct-mapped cache? Select all that apply.

Replacement: LRU

0%

Replacement: MRU

0%

Replacement: FIFO

0%

Replacement: Random

0%

Replacement: None of the above

0%

Write: Write-through

0%

Write: Write-back

0%

Direct Mapped: What is Possible?

- What policies can be implemented for a **direct-mapped cache**?
Select all that apply.

1. Block Replacement (Eviction) Policy

- A. Least Recently Used
- B. Most Recently Used
- C. FIFO
- D. Random
- ☒ E. None of the above

2. Write Policy

- ☒ A. Write-through
(memory access per write)
- ☒ B. Write-back
(dirty bit and write to memory on eviction)

In direct mapped caches, there is only ever one block to evict—the existing block with matching index.

Types of Misses

- Compulsory Miss
 - Caused by the first access to a block that has never been in the cache
- Capacity Miss
 - Caused when the cache cannot contain all the blocks needed during the execution of a program
 - Occur when blocks were in the cache, replaced, and later retrieved
- Conflict Miss
 - Multiple blocks compete for the same location in the cache, even when the cache has not reached full capacity.

Types of Misses

- **Compulsory Miss**
 - Caused by the first access to a block that has never been in the cache
- **Capacity Miss**
 - Caused when the cache cannot contain all the blocks needed during the execution of a program
 - Occur when blocks were in the cache, replaced, and later retrieved
- **Conflict Miss**
 - Multiple blocks compete for the same location in the cache, even when the cache has not reached full capacity.

Which types of misses can occur for Fully Associative caches (FA)?
For Direct Mapped caches (DM)? Select all that apply.

- A. Compulsory
- B. Capacity
- C. Conflict
- D. None of the above
- E. All of the above

Which types of misses can occur for Fully Associative caches (FA)? For Direct Mapped caches (DM)? Select all that apply.

FA: Compulsory

0%

FA: Capacity

0%

FA: Conflict

0%

FA: None of the above

0%

DM: Compulsory

0%

DM: Capacity

0%

DM: Conflict

0%

DM: None of the above

0%



Types of Misses

- Compulsory Miss
 - Caused by the first access to a block that has never been in the cache
- Capacity Miss
 - Caused when the cache cannot contain all the blocks needed during the execution of a program
 - Occur when blocks were in the cache, replaced, and later retrieved
- Conflict Miss
 - Multiple blocks compete for the same location in the cache, even when the cache has not reached full capacity.
 - Occurs in direct mapped caches, as well as in (out of scope) set associative caches.
 - Misses of this type would **not occur in a fully associative cache** with similar specs.

Summary: Cache Comparisons

■ Fully Associative

- A specific line of data can be stored in **any line** of the cache.
- No index.
- Need to choose replacement policy.
- Need to choose write policy.

■ Direct Mapped

- A specific line of data can only be stored in **one index** of the cache.
- Has index.
- If the line you want to store the data in is occupied, you kick out that line.
- Need to choose write policy.